

Optimization for Neural Networks

Aditya Ramesh

Abstract

TODO: Update this!

We introduce the concept of a neural network as a biologically-inspired graphical template that allows us to construct nonlinear functions to perform classification and regression tasks. The nonlinear function represented by the neural network is controlled by a set of parameters that are calibrated during the training process. By viewing the update rule for the optimization algorithm used during training as a discrete dynamical system, we can gain some insight into the stability properties of a neural network. Using this insight, we derive an optimal value for the learning rate used during training.

1 Introduction to Neural Networks

1.1 Representation

Neural networks take their inspiration from information processing in biological systems. In the human brain, information is transmitted between neurons in the form of electrical and chemical signals sent across synapses. Each neuron receives signals from a set of input neurons, and, under certain conditions, will broadcast a signal to a set of output neurons. This response can be thought of as the result of a local computation involving the input signals. We can model the flow of information in this network of neurons using a directed graph G , in which the neurons are nodes and the synapses are edges. This rough biological conceptualization of neural communication leads to a mathematical structure that we will fashion into a model of computation.

Our goal will now be to derive a mathematical description for G , so that the resulting neural network can be used for function approximation. To emphasize that our notion of neural network has little to do with neuroscience, we will refer to the “neurons” in the network as nodes in G . We limit our discussion to *feed-forward* neural networks, which do not contain cycles. Consequently, G must be a directed acyclic graph (DAG). Now suppose that we wish to use G to approximate a function $f : \mathbb{R}^m \rightarrow \mathbb{R}^n$. If $f(x) = y$ for some $x \in \mathbb{R}^m$ and $y \in \mathbb{R}^n$, then our task will be to use G to “learn” what f does to transform x into y .

Let us first impose some organizational structure on G . We delegate to a set I of m input nodes the task of broadcasting the components of x to other nodes in G . In order for the neural network to be useful, the transmission of information must eventually

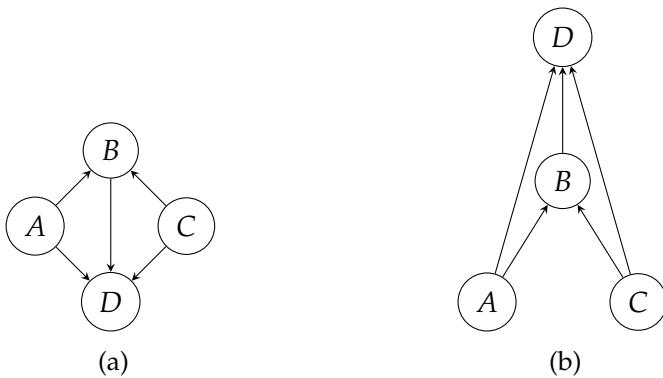


Figure 1.1: In 1.1a, we see a DAG, and in 1.1b, the representation of the DAG as a layered graph. Notice that the edges $A \rightarrow D$ and $C \rightarrow D$ are manifested as skip connections in the layered graph.

cease at a set O of n output nodes. The information computed by each node in O will correspond to a component of \hat{y} , the network's approximation to y . As things stand, I and O can consist of arbitrary nodes of G . It would help if we could hide the entangled mass of nodes of edges (aka *connections*) involved in the communication between I and O , and simply think of I as x and O as \hat{y} . Fortunately, our existing definitions allow us to do far more than this.

Any DAG can be rendered as a layered graph. In a layered graph, nodes are arranged in horizontal rows, and only vertical connections in one direction, between nodes in different layers, are allowed. The process by which this is accomplished is called *layered graph drawing* (see Figure 1.1). Since G is a DAG, we can partition its nodes into an array of successive layers L_0, \dots, L_d , where d is the *depth* of the neural network. The first layer, L_0 , consists of the input nodes ordered from left to right based on the components of x to which they correspond. It is called the *input layer*. Not all output nodes necessarily belong to the last layer, so an *output layer* need not exist in general. In our case, we assume it does, so L_d must be the output layer. Layers in between L_0 and L_d are called *hidden layers*. Figure 1.2 depicts a small, two-layer neural network.

Computation in a neural network proceeds from layer to layer. Nodes in a layer are called *units*, and edges between units (which must necessarily be from different layers) are called *connections*. As in Figure 1.1b, connections between units in nonconsecutive layers can occur; these are called *skip connections*. Consequently, in order for information to propagate from L_k to L_{k+1} , we may require the outputs of all units from L_0 to L_{k-1} . Let $w := |L_k|$ denote the *width* of L_k . Rather than thinking about L_k as a subgraph of G , we associate L_k with a vector $z_k \in \mathbb{R}^w$. The value of component z_{ki} is given by the output of the i th unit from the left end of L_k . It is now evident that the vector x is

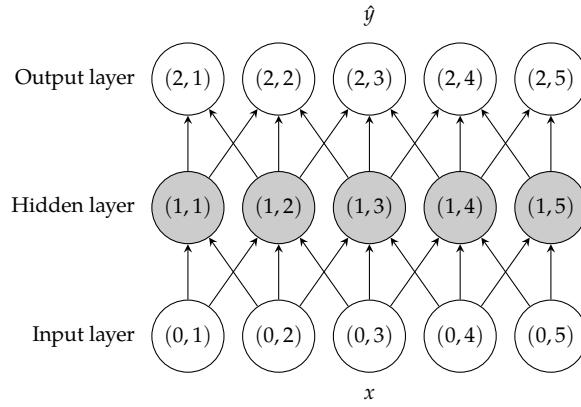


Figure 1.2: A small, two-layer neural network. The pair (k, j) is used to denote the j th unit in layer k .

transformed into \hat{y} by a series of successive functions $\sigma_1, \dots, \sigma_d$ given by

$$\begin{aligned}
x &=: z_0 \xrightarrow{\sigma_1} z_1 \\
(z_0, z_1) &\xrightarrow{\sigma_2} z_2 \\
(z_0, z_1, z_2) &\xrightarrow{\sigma_3} z_3 \\
&\vdots \quad \vdots \\
(z_0, \dots, z_{d-1}) &\xrightarrow{\sigma_d} z_d := \hat{y},
\end{aligned} \tag{1.1}$$

where σ_k is realized by the units in L_k .

The j th unit in the input layer of a neural network simply returns the corresponding component x_j of the input vector x . On the other hand, each unit in the hidden and output layers forms a linear combination of the inputs from previous layers to which it is connected, and adds a bias parameter to the result. Suppose that L_k is not the input layer, so that $k > 0$. The j th unit in L_k is denoted by (k, j) , and its activation a_{kj} is defined as

$$a_{kj} := \sum_{(l,i) \in \text{Pa}(k,j)} w_{li \rightarrow kj} z_{li} + b_{kj}. \tag{1.2}$$

Here, $\text{Pa}(k,j)$ denotes the parents of (k, j) , which are the units that have edges directed toward (k, j) . This notation allows us to easily generalize our discussion to layered networks that incorporate skip connections. The number $w_{li \rightarrow kj}$ is the *weight* associated with the connection $(l, i) \rightarrow (k, j)$, and b_{kj} the *bias* associated with (k, j) . The quantity z_{li} is the output of unit (l, i) , and is obtained by applying the unit's *activation function* $u_{li} : \mathbb{R} \rightarrow \mathbb{R}$ to the unit's activation a_{li} :

$$z_{li} = u_{li}(a_{li}). \tag{1.3}$$

Figure 1.3 shows how the activation of a unit with three parents is computed.

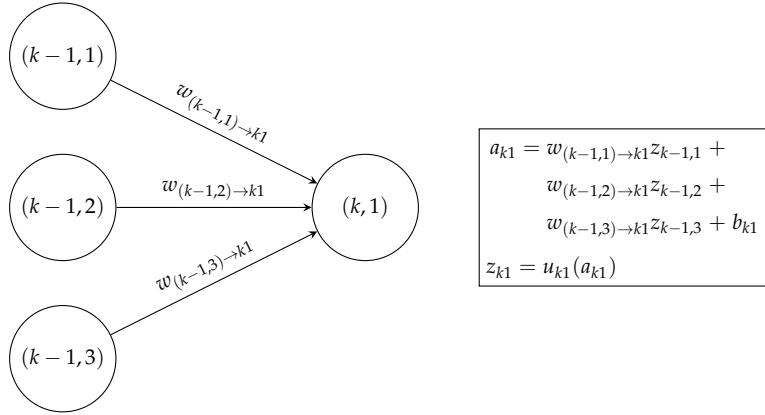


Figure 1.3: An illustration how the output of $(k, 1)$, a unit with three inputs from L_{k-1} , is computed. In this case, $\text{Pa}(k, 1) = \{(k-1, 1), (k-1, 2), (k-1, 3)\}$. First, the activation a_{k1} is computed by forming a linear combination of the outputs of the units in $\text{Pa}(k, 1)$ with the corresponding weights to (k, j) , and adding the bias b_{k1} . Then, the output z_{k1} is determined by applying the activation function u_{k1} to a_{k1} .

The process of computing the output vector z_k corresponds to applying the function σ_k . If $w := |L_k|$, then writing

$$(z_0, \dots, z_{k-1}) \xrightarrow{\sigma_k} z_k$$

is shorthand for saying that we compute z_k through w simultaneous applications of (1.2), followed by w simultaneous applications of (1.3). In order to compute the network's prediction $\hat{y} \in \mathbb{R}^m$ corresponding to an input $x \in \mathbb{R}^n$, we apply the functions $\sigma_1, \dots, \sigma_d$ as described by (1.1). This process is called *forward propagation*.

1.2 Activation Functions and Feature Learning

So far in our discussion, we have avoided making specific choices for the activation function $u_{kj} : \mathbb{R} \rightarrow \mathbb{R}$ associated with unit (k, j) . Three of the most commonly-used activation functions in the literature are the identity, scaled tanh, and linear threshold functions (see Figure 1.4). (In the neural network literature, linear threshold functions are called "rectified linear units" (ReLU), an instance of specialized terminology that Mehryar Mohri despises.) Typically, all units in the same layer are associated with the same activation function. We therefore drop the redundant second index in the subscript of u_{kj} , and simply refer to the activation function as u_k .

The choices for the constants a and b of the scaled tanh function are motivated by several reasons [LeCun et al., 1998]. Firstly, the scaled tanh function is symmetric and approximately linear about the origin. If the inputs are decorrelated, our choices for a and b cause the variance of the scaled tanh function to be close to unity. This accelerates the convergence of optimization algorithms (see Section 1.5).

Secondly, the choices for these constant helps prevent the *vanishing gradient problem*. Using these constants causes the second derivative of the scaled tanh function to be

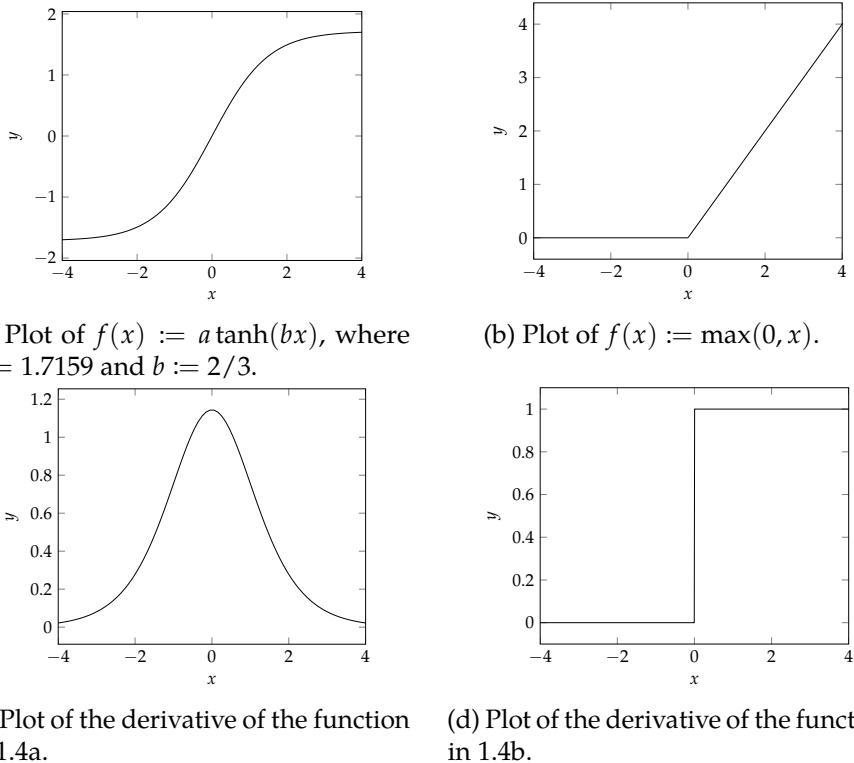


Figure 1.4: Plots of two of the most common activation functions and their derivatives.

maximized at ± 1 . This is intentional: ± 1 are the binary target values that are typically used for classification in practice. Suppose that the asymptotes of the scaled tanh function had been at the ± 1 instead of $\pm a$. Then the weights of the network would drastically increase, in order to produce the large activation values necessary to drive the outputs of tanh function to the target values at its asymptotes. The derivative of the scaled tanh function at these large activation values would be exponentially small (see Figure 1.4c). Thus, a minimization algorithm using derivative information would be liable to “get stuck”.

Finally, the choices for these constants allow us to interpret the output of the network as a measure of confidence towards its classification. When an input is near the decision boundary separating instances in two classes, we would like the network to output a small confidence value to reflect this uncertainty. If the asymptotes of the scaled tanh function were at ± 1 , then the large activation values would force the outputs to one of the two extreme values, regardless of the certainty of the classification. Our choices for a and b avoid this problem.

Unlike the scaled tanh function, the identity function is typically only used for regression. One common configuration for regression involves alternating between layers using the scaled tanh and identity activation functions. For classification, we typically only use the tanh or linear threshold functions. The adoption of the linear threshold function

in the literature is relatively recent [Nair and Hinton, 2010], but many have found that it drastically reduces training time and improves generalization when compared to the scaled tanh function [Krizhevsky et al., 2012]. Despite these general guidelines, choosing the best activation functions for a particular problem can involve a certain degree of experimentation.

One of the principles underlying the effectiveness of neural networks is that of *hierarchical feature learning*, and is the focus of much of the deep learning research at NYU. A *feature* is a “simple” quantity derived from the input that is designed to identify one or more of the input’s salient characteristics. For example, suppose that we wish to identify whether a given black and white image contains a human face. One potential feature is the difference between the sum of intensities of the pixels in the left half of the image and that of the right half of the image. In neural networks, a feature is a subset of activation values that become large when a given pattern is present in the input. This is a result of the weights and biases of the network being attuned to presence the pattern.

The success of neural networks in image recognition and acoustic modeling has been attributed to their ability to learn a *hierarchy* of features. In these applications, the location of a layer L_k in the network determines the relative scale of the patterns in the input captured by the features. In some sense, the lowest layers of the network “zoom in” to the input to capture low-level, local details, while the highest layers of the network “zoom out” to capture overarching, global patterns (see Figure 1.5). It is plausible that this phenomenon can occur, since each σ_k synthesizes the information from layers L_0, \dots, L_{k-1} to produce z_k .

The activation functions discussed so far are useful for learning the low- and mid-level features in the hierarchy. For learning high-level features, *radial basis functions* (RBFs) are sometimes more appropriate [LeCun et al., 2001]. A radial basis function $\phi : \mathbb{R} \rightarrow \mathbb{R}$ is a continuous function whose value depends only on the *radius* from a prescribed center $c \in \mathbb{R}^n$. Suppose that we wish to approximate a function $f : \mathbb{R}^n \rightarrow \mathbb{R}$ whose values at the points c_1, \dots, c_m are known. By centering an RBF at each c_i , we can approximate f using the function $s : \mathbb{R}^n \rightarrow \mathbb{R}$ given by

$$s(x) := \sum_{i=1}^m \lambda_i \phi(\|x - c_i\|_2), \quad (1.4)$$

where $\lambda \in \mathbb{R}^m$ is an adjustable parameter vector. The most commonly-used RBFs are the Euclidean and Gaussian basis functions. These functions are given by

$$r \mapsto r \quad \text{and} \quad r \mapsto \exp(-kr^2/2),$$

respectively, where $k \in \mathbb{R}^+$.

Since RBFs are used for learning high-level features, they are usually embedded into one of the topmost layers of the network. As an example, suppose that we wish to classify 32×32 black and white images of handwritten digits into ten classes, where each class corresponds to a number from zero to nine. Suppose that L_k is a layer with 32^2 units. By attaching a ten-unit layer L_{k+1} with RBF activation functions to L_k , we can generate ten



Figure 1.5: Visualization of the hierarchy of features in a fully-trained convolutional network for image classification, taken from Zeiler and Fergus [2014]. For each of the first five layers of the network, a selection of the highest activation values in the layer is shown alongside the corresponding input images. Note that in general, visualizing features in a meaningful way is very difficult.

scores from $z_k \in \mathbb{R}^{32 \times 32}$. The i th score is given by $(z_{k+1})_i$, and measures the confidence that the input $x \in \mathbb{R}^n$ belongs to the i th class.

The idea is to get z_k to match a *prototype* for one of the classes as closely as possible. We can view the functions $\sigma_1, \dots, \sigma_k$ as nonlinearly deconstructing and reassembling x to match one of these prototypes. The ten prototypes c_1, \dots, c_{10} , are initialized by running a clustering algorithm, such as k -means, on the training sample. Each prototype is treated as the center of an RBF, and is subsumed into the weights of the network. The proximity of z_k to the i th prototype c_i is given by $\phi(\|x - c_i\|)$, and is a measure of confidence that x belongs to the i th class.

If ϕ is a *normalized RBF*, then the confidence scores are probabilities. A normalized RBF ϕ is the normalized form of another basis function ψ . Hence, the i th RBF is given by

$$\phi(x) := \frac{\psi(\|x - c_i\|)}{\sum_{j=1}^{10} \psi(\|x - c_j\|)}. \quad (1.5)$$

Now the ten outputs, $\phi(\|x - c_i\|)$, $i \in [1, 10]$, form a discrete probability distribution over the classes.

1.3 Classification and Regression

We have seen that a neural network is a biologically-inspired nonlinear function controlled by two adjustable vectors of parameters: the weights w and the biases b . Now consider a sample $S := \{(x_1, y_1), \dots, (x_s, y_s)\}$, where each $x_k \in \mathbb{R}^n$ is an input vector, and each y_k is the target vector that we aim to predict when we are given x_k . We focus on two categories of tasks that we can perform using neural networks: classification and regression. For regression, we have $y_k \in \mathbb{R}^m$, so the output layer consists of a single unit.

For classification, two encoding schemes are possible. The first scheme can only be used for *unary classification*, in which we wish to classify x_k into one of several mutually-exclusive classes. In this scheme, each class is associated with an index, so $y_k, \hat{y}_k \in \mathbb{Z}$. These values are called *place codes*. A special case of unary classification is *binary classification*, in which we seek to place x_k into one of two mutually-exclusive classes. When using place codes, the output layer consists of a single unit.

The second scheme can be used for unary classification as well as *multiclass classification*, in which the x_k can belong to one or more classes. This time, $y_k \in \{0, 1\}^m$ and $\hat{y}_k \in \mathbb{R}^m$. These values are called *distributed codes*. The component $(y_k)_i$ is one if x_k belongs to the i th class, and zero otherwise. On the other hand, $(\hat{y}_k)_i$ is a measure of confidence that x_k belongs to the i th class. Whether a larger value indicate increased or decreased confidence varies based on convention.

Distributed codes possess the advantage that they scale well to large numbers of classes, while place codes do not [LeCun et al., 1998]. When using place codes, the single output unit of the network must assume one of finitely many values. If a large number of classes occur with nontrivial probabilities, then each unit in the penultimate layer of the network must generate an output close to zero almost all of the time. This becomes increasingly difficult as the number of classes grows.

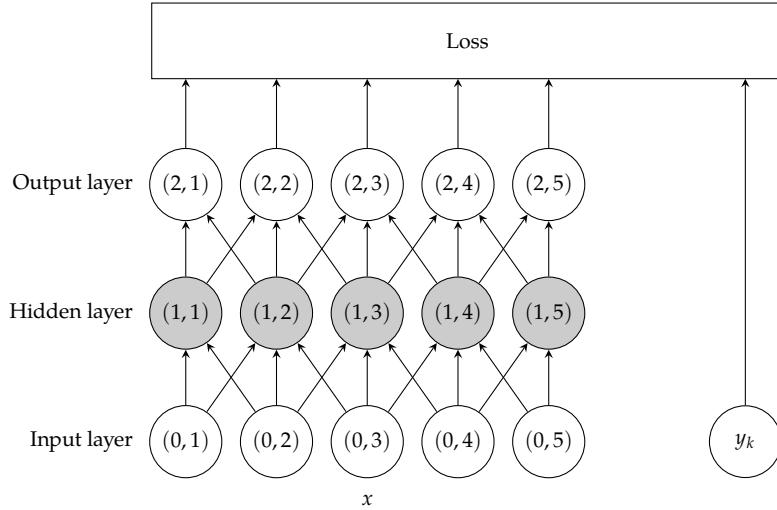


Figure 1.6: A visualization of how the k th instance (x_k, y_k) is fed into the two-layer neural network from Figure 1.2.

1.4 The Loss Function

Our goal is now to formulate the objective for the minimization algorithm used to adjust the weights and biases. To do this, we need a function that measures how well the network's prediction $\hat{y}_k \in \mathbb{R}^m$ for $x_k \in \mathbb{R}^n$ matches the target value $y_k \in \mathbb{R}^m$. This function is called the *loss function* or *cost function* (see Figure 1.6). The choice of the loss function depends on the type of problem we would like to solve.

One way to arrive at the loss function is to assign a probabilistic interpretation to the network outputs [Bishop, 2006]. Let w and b be the vectors of weights and biases of the network. By choosing an explicit representation for the conditional distribution $p(y_k | x_k, w, b)$, we are led to a canonical loss function corresponding to this choice. Let $X = \{x_1, \dots, x_s\}$ and $Y = \{y_1, \dots, y_s\}$. If we assume that the instances $(x_k, y_k) \in S$ are iid, then we can apply the chain rule to write

$$p(Y | X, w, b) = \prod_{k=1}^s p(y_k | x_k, w, b).$$

The LHS of (1.4) is called the likelihood function. The canonical loss function is obtained by maximizing the likelihood function over the w and b . Maximizing the likelihood function is equivalent to minimizing its negative logarithm, which is given by

$$-\ln p(Y | X, w, b) = -\sum_{k=1}^s \ln p(y_k | x_k, w, b). \quad (1.6)$$

The LHS of (1.6) is called the negative log-likelihood function (NLL).

For regression, it is often appropriate to assume that

$$p(y_k | x_k, w, b) = N(y_k | \hat{y}_k, \beta^{-1}),$$

where N is the normal distribution, which for mean μ and variance σ^2 is defined as

$$N(x | \mu, \sigma^2) := \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{1}{2\sigma^2}(x - \mu)^2\right).$$

The value \hat{y}_k is the network's prediction for x_k , and β is the precision (inverse variance) of the Gaussian noise. The corresponding NLL function is given by

$$\frac{\beta}{2} \sum_{k=1}^s (\hat{y}_k - y_k)^2 - \frac{s}{2} \ln \beta + \frac{s}{2} \ln(2\pi).$$

We discard the constant terms, which do not affect the minimization process. This leaves us with the sum-of-squares loss function, which is given by

$$E(w, b) := \frac{1}{2} \sum_{k=1}^s (\hat{y}_k - y_k)^2 = \sum_{k=1}^s e(\hat{y}_k, y_k), \quad (1.7)$$

where e is the per-instance loss function

$$e(\hat{y}_k, y_k) := \frac{1}{2} (\hat{y}_k - y_k)^2.$$

The sum-of-squares error function is generally inappropriate for classification. We adopt the use of distributed codes, so $y_k \in \{0, 1\}^m$ and $\hat{y}_k \in \mathbb{R}^m$. We further suppose that the components of \hat{y}_k determine a discrete probability distribution over the classes. For convenience, we define $y_{ki} := (y_k)_i$ and $\hat{y}_{ki} := (\hat{y}_k)_i$. Our assumption is that the conditional distribution is categorical, and given by

$$p(y_k | x_k, w, b) = \prod_{i=1}^m (\hat{y}_{ki})^{y_{ki}}.$$

The corresponding NLL function is the *multiclass cross-entropy function*, and does not have any constant terms that we can discard. So we define

$$E(w, b) = - \sum_{k=1}^s \sum_{i=1}^m y_{ki} \ln \hat{y}_{ki} = \sum_{k=1}^s e(\hat{y}_k, y_k),$$

where e is the per-instance loss function

$$e(\hat{y}_k, y_k) := - \sum_{i=1}^m y_{ki} \ln \hat{y}_{ki}.$$

In the special case where we are performing binary classification, with $y_k \in \{0, 1\}$ and $\hat{y}_k \in \mathbb{R}$, we can write

$$p(y_k | x_k, w, b) = (\hat{y}_k)^{y_k} (1 - \hat{y}_k)^{1-y_k}$$

and

$$e(\hat{y}_k, y_k) := -(y_k \ln \hat{y}_k + (1 - y_k) \ln(1 - \hat{y}_k)).$$

The cross-entropy loss function is almost exclusively used with a special activation function for the output layer of the network. This activation function is the *softmax* function; for $z \in \mathbb{R}^m$, it is given by

$$\phi(z_i) = \frac{\exp(z_i)}{\sum_{j=1}^m \exp(z_j)}. \quad (1.8)$$

This is a normalized RBF of the form given by (1.5). It follows that using (1.8) as the output activation function will guarantee that the components of \hat{y}_k determine a discrete probability distribution over the classes. For binary classification, (1.8) reduces to the *logistic sigmoid* function upon setting the component \hat{y}_{k2} , which corresponds to the second class, to zero:

$$\begin{aligned}\phi(\hat{y}_1) &= \frac{\exp(\hat{y}_1)}{\exp(\hat{y}_1) + \exp(\hat{y}_2)} \\ &= \frac{\exp(\hat{y}_1)}{\exp(\hat{y}_1) + 1} \\ &= \frac{1}{1 + \exp(-\hat{y}_1)}.\end{aligned}$$

Thus, the softmax function can be viewed as a generalization of the logistic sigmoid function to more than two classes.

1.5 Data Preprocessing and Initialization

Preprocessing the data can have drastic effects on convergence of the optimization algorithm used for learning. However, normalizing the features (the components of an input $x_k \in \mathbb{R}^n$) must be done with great care, so as not to destroy useful information encoded by the relative magnitudes of a feature. Suppose that it is equally important that we predict the target value for any input x_k from our training sample D . This is *not* the case, for instance, when we wish to predict whether the person corresponding to an input x_k has a rare disease. Then it can be beneficial to center and scale the features by subtracting the mean vector μ and scaling by the variance vector σ . Specifically, each input x_k is normalized by computing

$$\bar{x}_k := \frac{x_k - \mu}{\sigma}. \quad (1.9)$$

In general, it is recommended to *decorrelate* the features whenever possible [LeCun et al., 1998]. Doing so decouples the updates to the individual weights of the network during the optimization process, which accelerates convergence. LeCun et al. [1998] give an intuitive explanation for why decorrelating the inputs can be beneficial. Consider the extreme case where all components of the inputs are positive. Let $(1, j)$ be a unit in the second layer of the network, and let \tilde{x}_k consist of the components of x_k corresponding to $\text{Pa}(1, j)$ (see Figure 1.2). Let δ be the scalar error computed for $(1, j)$ using the techniques

discussed in Section 1.6. Then a naïve first-order optimization algorithm will update the weights to $(1, j)$ by an amount proportional to $\delta \tilde{x}_k$. Thus, the weights to $(1, j)$ will either *all* increase or *all* decrease (depending on the sign(δ)) after each update. This makes finding a good weight configuration difficult. A more sophisticated approach to decorrelating the inputs involves PCA [LeCun et al., 1998, Krizhevsky et al., 2012].

The initial values of the weights can have a significant impact on both convergence speed and generalization error. Suppose that we use the scaled tanh activation function. In order to avoid vanishing gradient problem, we would like the activation values to remain in the range $[-1, 1]$, which is within the linear region of the scaled tanh function. One way of enforcing this is to require that the activation values have zero mean and unit variance [LeCun et al., 1998].

Suppose that the inputs are centered and scaled according to (1.9), so that

$$\mathbb{E}(x_i) = 0 \quad \text{and} \quad \text{Var}(x_i) = 1,$$

for each $i \in [1, n]$. From (1.2), we know that the activation of $(1, j)$ is given by

$$a_{1j} := \sum_{(0,i) \in \text{Pa}(1,j)} w_{0l \rightarrow 1j} x_{0l} + b_{1j}.$$

If the weights are chosen independently from the inputs, and both the weights and the biases have zero mean, then we have

$$\begin{aligned} \mathbb{E}(a_{1j}) &= \mathbb{E}\left(\sum_{(0,i) \in \text{Pa}(1,j)} w_{0l \rightarrow 1j} x_{0l} + b_{1j}\right) \\ &= \sum_{(0,i) \in \text{Pa}(1,j)} \mathbb{E}(w_{0l \rightarrow 1j} x_{0l} + b_{1j}) \\ &= \sum_{(0,i) \in \text{Pa}(1,j)} (\mathbb{E}(w_{0l \rightarrow 1j} x_{0l}) + \mathbb{E}(b_{1j})) \\ &= \sum_{(0,i) \in \text{Pa}(1,j)} \mathbb{E}(w_{0l \rightarrow 1j}) \mathbb{E}(x_{0l}) = 0. \end{aligned}$$

Thus, our requirement that the activation values have zero mean is satisfied.

We now describe how to enforce that $\text{Var}(a_{1j}) = 1$. Let $f := |\text{Pa}(1,j)|$, and suppose that all weights have the same standard deviation σ . The value f is called the *fan-in* of

$(1, j)$. To solve for the value of σ that forces $\text{Var}(a_{1j}) = 1$, we compute

$$\begin{aligned}
\text{Var}(a_{1j}) &= \text{Var} \left(\sum_{(0,i) \in \text{Pa}(1,j)} w_{0l \rightarrow 1j} x_{0l} + b_{1j} \right) \\
&= \text{Var} \left(\sum_{(0,i) \in \text{Pa}(1,j)} w_{0l \rightarrow 1j} x_{0l} + \sum_{(0,i) \in \text{Pa}(1,j)} b_{1j} \right) \\
&= \text{Var} \left(\sum_{(0,i) \in \text{Pa}(1,j)} w_{0l \rightarrow 1j} x_{0l} \right) \\
&= \sum_{(0,i) \in \text{Pa}(1,j)} (\text{Var}(w_{0l \rightarrow 1j}) \text{Var}(x_{0l}) - \text{E}(w_{0l \rightarrow 1j})^2 \text{E}(x_{0l})^2) \\
&= \sum_{(0,i) \in \text{Pa}(1,j)} \text{Var}(w_{0l \rightarrow 1j}) \text{Var}(x_{0l}) \\
&= \sum_{(0,i) \in \text{Pa}(1,j)} \sigma^2 = f\sigma^2.
\end{aligned}$$

So in order to enforce that $\text{Var}(a_{1j}) = 1$, we should sample the weights from a distribution with zero mean and standard deviation given by

$$\sigma = \frac{1}{\sqrt{f}}.$$

The most common choices of distribution for weight initialization are the uniform and normal distributions [LeCun et al., 1998, Krizhevsky et al., 2012].

The training dynamics of neural networks are vitally important but poorly understood. As the weights and biases of the network are calibrated using the minimization algorithm, the distributions of the activation values at each layer evolve. For deep networks with many layers, the vanishing gradient problem is greatly compounded. In particular, the greater the depth of a layer in the network, the greater the tendency of the activation values to cluster around zero or become very large [Glorot and Bengio, 2010]. This causes the vanishing gradient problem discussed in Section 1.2. The subtle numerical issues resulting from these complex interactions largely determines whether, when, and how a given minimization algorithm will converge.

The use of a good weight initialization scheme can help mitigate the undesirable saturation of the activation values [Glorot and Bengio, 2010]. Several authors have recently proposed new initialization schemes that they found were more effective than the one given here [Glorot and Bengio, 2010, Martens, 2010]. Another strategy is to incorporate a form of internal normalization into the network by adding *local response normalization* layers [Krizhevsky et al., 2012]. Part of the reason that training deep neural networks is so difficult is that these dynamics are so poorly understood. Development of a deeper understanding of what is happening in these situations is likely to have widespread practical consequences.

1.6 Computing Derivatives

TODO

1.7 Convolutional Networks

TODO

2 Concepts from Statistical Learning Theory

2.1 Generalization and PAC-learning

Our initial goal when motivating the discussion for neural networks was to “learn” what a given function $f : \mathbb{R}^n \rightarrow \mathbb{R}^m$ does to transform $x \in \mathbb{R}^n$ to $y \in \mathbb{R}^m$. But what exactly does it mean to learn? So far, we have neglected to provide a definition for what it means for an algorithm “learn” a model to solve a given problem. A rigorous definition for learning is given in the field of statistical learning theory. This definition will have widespread consequences in shaping our criteria used to determine when an optimization algorithm is good.

When performing binary classification, each training instance $(x_k, y_k) \in S$ is implicitly associated with a target *concept* that we wish to learn. This concept is a mapping c from the input space X to the output space $Y := \{0, 1\}$, such that $y_k = c(x_k)$. The set of all concepts associated with our classification task is called the *concept class*, and is denoted by C . A concept can be something simple, such as the set of axis-aligned rectangles in \mathbb{R}^2 , or something exceedingly complex, such as the manifold of human faces in $\mathbb{R}^{256 \times 256}$, the space of 256×256 black and white images. We assume that all instances are drawn iid from some unknown distribution D over X .

Suppose that we are given a family of candidate functions F , such that $f : X \rightarrow \{0, 1\}$ for every $f \in F$. The task of our learning algorithm is to minimize the probability of a misclassification from an instance $x \sim D$. This leads to the following definition for *generalization error*.

Definition 1 (Generalization error [Mohri et al., 2012]). *Given a candidate function $f \in F$, a target concept $c \in C$, and an underlying distribution D , the generalization error or risk of f is defined as*

$$R(f) := \Pr_{x \sim D}(f(x) \neq c(x)) = \mathbb{E}_{x \sim D}(1\{f(x) \neq c(x)\}).$$

The generalization error is not a quantity that we can directly measure. Instead, we must work with the *empirical error*, the average number of misclassifications over our training sample $S \sim D$.

Definition 2 (Empirical error [Mohri et al., 2012]). *Given a candidate function $f \in F$, a target concept $c \in C$, and a sample $S := \{(x_1, y_1), \dots, (x_s, y_s)\} \sim D$, the empirical error or*

empirical risk of f is defined as

$$\hat{R}(f) := \frac{1}{s} \sum_{i=1}^s \mathbb{1}\{f(x_i) \neq c(x_i)\}.$$

There is a simple relationship between the empirical and generalization errors. Let $f \in F$. Using the linearity of expectations and the fact that S is drawn iid from D , it is easy to show that

$$\mathbb{E}_{S \sim D}(\hat{R}(f)) = R(f).$$

In other words, the generalization error of f is simply the average misclassification rate over samples drawn from D . We are now ready to present the definition of probably approximately correct (PAC) learning.

Definition 3 (PAC-learning [Mohri et al., 2012]). *Suppose that we are performing binary classification of inputs from an input space X , and let C be the associated concept class. Let $O(n)$ be an upper bound on the space complexity of an input $x \in X$. Given an algorithm A and a sample $S \sim D$, let $f_S \in F$ denote the candidate function found by A for S . Then C is PAC-learnable if there exists an algorithm A and a polynomial function $g : \mathbb{R}^4 \rightarrow \mathbb{R}$ such that for any $\epsilon, \delta \in (0, 1]$, for all distributions D over X , and for any target concept $c \in C$, the following holds for any sample $S \sim D$ of size $s \geq g(1/\epsilon, 1/\delta, n, \text{size}(c))$:*

$$\Pr_{S \sim D}(R(f_S) \leq \epsilon) \geq 1 - \delta.$$

If A runs in $g(1/\epsilon, 1/\delta, n, \text{size}(c))$, then C is said to be efficiently PAC-learnable. In this case, A is called a PAC-learning algorithm for C .

The definition for PAC-learning is verbose, but the underlying intuition is simple. Suppose that we have a target generalization error ϵ that we wish to attain with probability $1 - \delta$ on any sample $S \sim D$. Then C is PAC-learnable if the sample size required to meet our goals is bounded by a polynomial function of $1/\epsilon$ and $1/\delta$. Roughly speaking, A must find a function in F that is *approximately correct with high probability*. But it must be able to do this using a sample size that is not “too large”. The extension of space and time complexity to sample size requires this notion of “high probability”.

One of the goals of statistical learning theory is to derive upper bounds on the generalization error in various situations. When such an upper bound applies to the particular situation at hand, we know that it is possible, at least in principle, to learn efficiently. Many of these upper bounds take the following form. Let $\delta \in (0, 1]$ be our target confidence, and let s be the size of the training samples drawn from D . Then with probability $1 - \delta$, it holds for any $f \in F$ that

$$R(f) \leq \hat{R}(f) + c \sqrt{\frac{\text{capacity}(F)}{s}}, \quad (2.1)$$

where $c \in \mathbb{R}^+$ and $\text{capacity}(F)$ is a quantity that measures the complexity of our chosen family of functions. When we apply the bound (2.1), we often subsume auxiliary constants into c . So two occurrences of c do not necessarily represent the same constant value.

2.2 ERM vs SRM

The central idea of statistical learning theory is that learning is a tradeoff between minimizing the training error and the complexity of the family of functions used to model C . These quantities correspond to the first and second terms on the RHS of (2.1), respectively. This tradeoff can be expressed in another way that is particularly illuminating. Let

$$f^* := \arg \min_f R(f) \quad \text{and} \quad f_F^* := \arg \min_{f \in F} R(f).$$

In other words, $R(f^*)$ is the best we can ever hope to do, while $R(f_F^*)$ is the best we can do using our chosen family F . Given a function $f \in F$, we define the *excess error* \mathcal{E} as

$$\mathcal{E} := R(f) - R(f^*) = (R(f_F^*) - R(f^*)) + (R(f) - R(f_F^*)) = \mathcal{E}_{\text{app}} + \mathcal{E}_{\text{est}}, \quad (2.2)$$

where

$$\mathcal{E}_{\text{app}} = R(f_F^*) - R(f^*) \quad \text{and} \quad \mathcal{E}_{\text{est}} = R(f) - R(f_F^*)$$

The quantity \mathcal{E}_{app} is called the *approximation error*, and measures how well functions in F can approximate functions in C . It is generally intractible, as we have no way of estimating it. The quantity \mathcal{E}_{est} is called the *estimation error*, and measures the performance of f relative to that of the best-in-class hypothesis f_F^* . It is easy to show that \mathcal{E}_{est} is bounded above by the second term on the RHS of (2.1). Equation 2.2 tells us that the quality of a candidate function f is determined by the sum of these errors.

Reducing the training error causes the estimation error to decrease, while reducing the complexity term causes the approximation error to decrease. But reducing the approximation error often requires increasing capacity(F), and this increases the complexity term. On the other hand, reducing the complexity term requires decreasing capacity(F), inhibiting our ability to approximate C . This increases the training error. The interplay between these two errors is called the *approximation-estimation tradeoff*.

Minimizing only the training error while ignoring the complexity term is called *empirical risk minimization* (ERM), and is inadvisable on both theoretical and practical grounds. An alternative procedure that controls capacity(F) is called *structural risk minimization* (SRM). Let $c \in \mathbb{R}^+$, and define

$$F_\lambda := \{f \in F : \text{complexity}(f) \leq c\}, \quad (2.3)$$

where $\text{complexity}(f)$ is a measure of complexity of f . Consequently, $F_a \subseteq F_b$ whenever $a, b \in \mathbb{R}^+$ such that $a < b$. ERM and SRM seek to find the functions given by

$$f_{\text{ERM}} := \arg \min_{f \in F} \hat{R}(f) \quad \text{and} \quad f_{\text{SRM}} := \arg \min_{\substack{c \in \mathbb{R}^+ \\ f \in F_c}} \hat{R}(f), \quad (2.4)$$

respectively. The SRM procedure finds a function in F that minimizes the generalization bound (2.1) (see Figure 2.1). We can transform second equation in (2.4) into an unconstrained minimization problem by using the penalty method. This gives

$$f_{\text{SRM}} = \arg \min_{\substack{\lambda \in \mathbb{R}^+ \\ f \in F}} (\hat{R}(f) + \lambda \text{complexity}(f)). \quad (2.5)$$

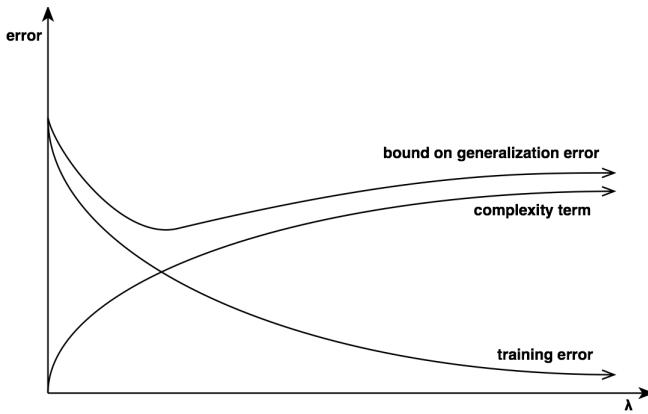


Figure 2.1: Visualization of the SRM procedure (adapted from Mohri et al. [2012]). The SRM procedure minimizes an upper bound of the generalization error that is a sum of the training error and the complexity term.

The Lagrange multiplier λ called the the *regularization parameter*, and penalizes more complex functions. We can now conduct our search over λ instead of over c . If H is a vector space, then we can define $\text{complexity}(f) := \|f\|$ for some norm $\|\cdot\|$. Note that SRM involves finding the solution to several ERM subproblems, and is intractable in general. In practice, we find an approximate SRM solution by using cross-validation to choose λ^* from a finite sequence $\{\lambda_k\}_{k=1}^r$.

As an example, suppose that we are performing least-squares regression, where $X = \mathbb{R}^n$ and $Y \in \mathbb{R}$. Let S be a training sample consisting of s instances. We define

$$F := \{x \mapsto w^t x + b : w \in \mathbb{R}^n, b \in \mathbb{R}\},$$

and assume let E be the sum-of-squares error given by (1.7). According to the theory developed so far, choosing $f := \arg \min_{f \in F} E(f)$ is highly inadvisable. Let $\text{complexity}(f) := \|w\|_2$ for any $f \in F$. A better idea would be to try sequence of values for λ in (2.5), such as $\{\lambda_k\}_{k=-5}^5$, where $\lambda_k := 2^k$. The results from statistical learning theory shape our understanding of how we should solve problems in machine learning.

3 Characterizing Optimization Algorithms

3.1 Goals of Optimization

We turn our focus to applying statistical learning theory to characterize the quality of optimization algorithms for machine learning. Let f_{ERM} be given by (2.4). During training, we only require that the optimization algorithm produce a function $\hat{f} \in F$, whose empirical error is within a fixed tolerance $\rho \in \mathbb{R}^+$ of that of f_{ERM} :

$$\hat{R}(\hat{f}) - \hat{R}(f_{\text{ERM}}) \leq \rho. \quad (3.1)$$

	capacity(F)	s	ρ
\mathcal{E}_{app}	\downarrow	—	—
\mathcal{E}_{est}	\updownarrow	\updownarrow	—
\mathcal{E}_{opt}	\uparrow	\uparrow	\downarrow
t	\uparrow	\uparrow	\downarrow

Table 3.1: The effects of increasing capacity(F), s , and ρ on the three errors and computing time. The symbol \updownarrow indicates that either an increase or a decrease is possible.

The expectation of the LHS of (3.1) taken over a sample $S \sim D$ is called the *optimization error* [Bousquet and Bottou, 2008], and is defined as

$$\mathcal{E}_{\text{opt}} := R(\hat{f}) - R(f_{\text{ERM}}).$$

We can now express the excess error of \hat{f} as

$$\mathcal{E} = (R(f_F^*) - R(f^*)) + (R(f_{\text{ERM}}) - R(f_F^*)) + (R(\hat{f}) - R(f_{\text{ERM}})) \quad (3.2)$$

$$= \mathcal{E}_{\text{app}} + \mathcal{E}_{\text{est}} + \mathcal{E}_{\text{opt}}. \quad (3.3)$$

In practice, the quality of the solution to a machine learning problem is determined by monetary cost (which determines the maximum sample size s_{\max}) and human patience (which determines the limit on computation time t_{\max}). The definition for PAC-learning is one way of characterizing when one might be satisfied with the performance of a learning algorithm with respect to these constraints. Let $t(F, \rho, s)$ represent the time required by the optimization algorithm to find a function $\hat{f} \in F$ satisfying $\mathcal{E}_{\text{opt}} \leq \rho$, given a sample of size s . Minimizing the excess error in (3.3) subject to the constraints on s_{\max} and t_{\max} yields the following meta-optimization problem [Bousquet and Bottou, 2008]:

$$\begin{aligned} & \underset{F, \rho, s}{\text{minimize}} \quad \mathcal{E} = \mathcal{E}_{\text{app}} + \mathcal{E}_{\text{est}} + \mathcal{E}_{\text{opt}} \\ & \text{subject to} \quad s \leq s_{\max} \\ & \quad t(F, \rho, s) \leq t_{\max}. \end{aligned} \quad (3.4)$$

Table 3.1 shows how the errors and computing time are affected when capacity(F), ρ , and s are increased.

As stated in Bousquet and Bottou [2008], “the solution of the optimization program (3.4) depends critically on which budget constraint is active: [the] constraint $s < s_{\max}$ on the number of examples, or [the] constraint $t < t_{\max}$ on the training time”. A *small-scale* problem is one that is constrained by the sample size. In this case, the computation time is not limited, so we can make \mathcal{E}_{opt} negligible by choosing ρ sufficiently small. Using the fact that \mathcal{E}_{est} is bounded above by the second term of (2.1), we have

$$\mathcal{E} = \mathcal{E}_{\text{app}} + \mathcal{E}_{\text{est}} \leq \mathcal{E}_{\text{app}} + c \sqrt{\frac{\text{capacity}(F)}{s}}. \quad (3.5)$$

A *large-scale* problem is one that is constrained by the maximum computing time. In this case, performing approximate optimization by choosing $\rho > 0$ improves generalization, because we can process more training instances in the allotted time. Using (2.1), it is easy to show that

$$\mathcal{E} = \mathcal{E}_{\text{app}} + \mathcal{E}_{\text{est}} + \mathcal{E}_{\text{opt}} \leq \mathcal{E}_{\text{app}} + \rho + c \sqrt{\frac{\text{capacity}(F)}{s}}. \quad (3.6)$$

Unfortunately, the bounds (3.5) and (3.6) are too pessimistic, and do not accurately reflect reality. In order to proceed with our analysis, we make some simplifying assumptions. First, we assume that F is the family of affine functions given by (2.2). In this special case, the generalization bound (2.1) simplifies to

$$R(f) \leq \hat{R}(f) + c \sqrt{\frac{n}{s}}.$$

A much sharper bound on the excess error can now be applied [Bousquet and Bottou, 2008]. It is given by

$$\mathcal{E} = \mathcal{E}_{\text{app}} + \mathcal{E}_{\text{est}} + \mathcal{E}_{\text{opt}} \leq c \left(\mathcal{E}_{\text{app}} + \left(\frac{n}{s} \log \frac{s}{n} \right)^\alpha + \rho \right), \quad (3.7)$$

where $\alpha \in [1/2, 1]$ is a constant that depends on a variance bound on the loss function. Since the three components of the excess error should decrease at approximately the same rate [Bousquet and Bottou, 2008], we make the additional assumption that

$$\mathcal{E} \approx \mathcal{E}_{\text{app}} \approx \mathcal{E}_{\text{est}} \approx \mathcal{E}_{\text{opt}} \approx \left(\frac{n}{s} \log \frac{s}{n} \right)^\alpha \approx \rho. \quad (3.8)$$

3.2 Analysis of Optimization Algorithms

We now investigate the properties of four specific optimization algorithms. Assume that our empirical loss function E is convex and twice differentiable (e.g. the sum-of-squares loss function). For convenience, we let $\theta \in \mathbb{R}^{2n}$, where the first n components of θ refer to a weight vector $w \in \mathbb{R}^n$ and the last n components of θ refer to a bias vector $b \in \mathbb{R}^n$. Thus $\theta = (w, b)$, for some $w, b \in \mathbb{R}^n$. We define $E(\theta) := E(w, b)$. Since E is convex, it has a unique minimum in terms of θ . Let $\theta_{\text{ERM}} := (w_{\text{ERM}}, b_{\text{ERM}})$ be the parameter vector corresponding to f_{ERM} . For a given input $x_k \in \mathbb{R}^n$, we define $\hat{y}_k^{\text{ERM}} := w_{\text{ERM}} x_k + b_{\text{ERM}}$.

The Hessian matrix H_{ERM} and gradient covariance matrix G_{ERM} evaluated at θ_{ERM} play an important role in our analyses [Bousquet and Bottou, 2008]. They are defined as

$$H_{\text{ERM}} := (\nabla_\theta^2 E)(\theta_{\text{ERM}}) = \frac{1}{s} \sum_{i=1}^s (\nabla_\theta^2 e)(\hat{y}_k^{\text{ERM}}, y_k)$$

and

$$\begin{aligned} G_{\text{ERM}} &:= ((\nabla_\theta E)(\theta_{\text{ERM}})) ((\nabla_\theta E)(\theta_{\text{ERM}}))^t \\ &= \frac{1}{s} \sum_{i=1}^s ((\nabla_\theta e)(\hat{y}_k^{\text{ERM}}, y_k)) ((\nabla_\theta e)(\hat{y}_k^{\text{ERM}}, y_k))^t, \end{aligned}$$

where e is the per-instance loss function associated with E . We assume that there exist constants $\lambda_{\max} \geq \lambda_{\min} > 0$ and $\nu > 0$ such that for any $\delta \in (0, 1]$, there exists a sample size $s \in \mathbb{N}$ such that, with probability at least $1 - \delta$, we have

$$\text{tr}(G_{\text{ERM}}(H_{\text{ERM}})^{-1}) \leq \nu \quad \text{and} \quad \text{Sp}(H_{\text{ERM}}) \subset [\lambda_{\min}, \lambda_{\max}].$$

The condition number $\kappa := \lambda_{\max}/\lambda_{\min}$ determines the difficulty of the optimization problem.

The optimization algorithms used in machine learning either operate on *full batches* or *mini-batches*. A full batch algorithm will only apply an update to θ after processing the entire training sample S . The derivative information computed by batch algorithms is averaged over the batch of instances. A mini-batch algorithm applies an update to θ every r iterations, where $r|s$. A *stochastic* or *online* algorithm uses a mini-batch size of one. The first two algorithms that we present are full batch algorithms, while the last two algorithms are stochastic.

Mini-batch algorithms are often several times faster than full batch algorithms. This is because mini-batch algorithms are better able to exploit redundancy in the training sample S . LeCun et al. [1998] give an explanation for why this happens. Suppose that $|S| = 1000$, and consists of the same 100 instances replicated ten times in the same order. The average of the gradient over S is equal to the average of the gradient over the first 100 instances. A full batch algorithm will compute the same gradient ten times before applying a single update. In this time, an algorithm with a mini-batch size of 100 would apply ten updates. While duplicate instances are rare in practice, there are often clusters of instances that are similar in some sense.

Following the precedent of Stephen Wright, we refer to the algorithms below as *gradient update* algorithms rather than as *gradient descent* algorithms, because the search directions used by the update rules are not necessarily descent directions.

- The gradient update (GU) algorithm uses the update rule

$$\theta_{k+1} := \theta_k - \eta (\nabla_{\theta} E)(\theta_k) = \theta_k - \frac{\eta}{s} \sum_{i=1}^s (\nabla_{\theta} e)(\hat{y}_k^{\text{ERM}}, y_k)$$

When $\eta = 1/\lambda_{\max}$, this algorithm requires $O(\kappa \log(1/\rho))$ updates to reach accuracy ρ .

- The second-order gradient update (2GU) algorithm uses the update rule

$$\theta_{k+1} := \theta_k - (H_{\text{ERM}})^{-1} (\nabla_{\theta} E)(\theta_k) = \theta_k - \frac{(H_{\text{ERM}})^{-1}}{s} \sum_{i=1}^s (\nabla_{\theta} e)(\hat{y}_k^{\text{ERM}}, y_k)$$

Note that the Hessian is *not* evaluated at each θ_k : we assume that we know H_{ERM} in advance, and use the same matrix at each iteration. In general, this algorithm requires $O(\log \log(1/\rho))$ iterations to reach accuracy ρ .

Algorithm	Cost of one iteration	Iterations to reach ρ	Time to reach accuracy ρ	Time to reach $\mathcal{E} \leq c(\mathcal{E}_{\text{app}} + \epsilon)$
GU	$O(ns)$	$O\left(\kappa \log \frac{1}{\rho}\right)$	$O\left(ns\kappa \log \frac{1}{\rho}\right)$	$O\left(\frac{n^2\kappa}{\epsilon^{1/\alpha}} \log^2 \frac{1}{\epsilon}\right)$
2GU	$O(n^2 + ns)$	$O\left(\log \log \frac{1}{\rho}\right)$	$O\left((n^2 + ns) \log \log \frac{1}{\rho}\right)$	$O\left(\frac{n^2}{\epsilon^{1/\alpha}} \log \frac{1}{\epsilon} \log \log \frac{1}{\epsilon}\right)$
SGU	$O(n)$	$\frac{\nu\kappa^2}{\rho} + o\left(\frac{1}{\rho}\right)$	$O\left(\frac{nv\kappa^2}{\rho}\right)$	$O\left(\frac{nv\kappa^2}{\epsilon}\right)$
2SGU	$O(n^2)$	$\frac{\nu}{\rho} + o\left(\frac{1}{\rho}\right)$	$O\left(\frac{nv}{\rho}\right)$	$O\left(\frac{nv}{\epsilon}\right)$

Table 3.2: Asymptotic rates for the four algorithms discussed (adapted from Bousquet and Bottou [2008]).

- Given an instance $(x_k, y_k) \in S$, the stochastic gradient update (SGU) algorithm uses the update rule

$$\theta_{k+1} := \theta_k - \frac{\eta}{k} (\nabla_{\theta} e)(\hat{y}_k^{\text{ERM}}, y_k).$$

If $\eta = 1/\lambda_{\min}$, this algorithm requires $\nu\kappa^2/\rho + o(1/\rho)$ iterations, on average, to reach accuracy ρ .

- Given an instance $(x_k, y_k) \in S$, the second-order stochastic gradient update algorithm (2SGU) uses the update rule

$$\theta_{k+1} := \theta_k - \frac{(H_{\text{ERM}})^{-1}}{k} (\nabla_{\theta} e)(\hat{y}_k^{\text{ERM}}, y_k).$$

This algorithm requires $\eta/\rho + o(1/\rho)$ iterations, on average, to reach accuracy ρ . Note that, unlike 2GU, the use of second-order information does not weaken the dependence on ρ .

The asymptotic rates for the four algorithms that were presented are given in Table 3.2. The rates in the first column are based on the sizes of the vectors and matrices involved in the corresponding update rules. The rates in the second column are taken directly from our discussion. To obtain the rates in the third column, we multiply the corresponding rates in the first and second columns. One of our assumptions in (3.8) is that

$$\rho \approx \left(\frac{n}{s} \log \frac{s}{n}\right)^{\alpha}. \quad (3.9)$$

Substituting this value of ρ into the bound (3.7) and comparing the result to $c(\mathcal{E}_{\text{app}} + \epsilon)$ shows that

$$\rho \approx \epsilon. \quad (3.10)$$

Using (3.9) and (3.10), it can then be shown that

$$s \approx \frac{n}{\epsilon^{1/\alpha}} \log \frac{1}{\epsilon}. \quad (3.11)$$

Substituting (3.10) and (3.11) into the rates in the third column gives the rates in the fourth column.

Table 3.2 contains a wealth of information regarding the properties of the four optimization algorithms. This information leads to the following conclusions for large-scale problems [Bousquet and Bottou, 2008]:

- Stochastic algorithms are more appropriate when there is less need to optimize accurately. This is the case in large-scale problems, or when the estimation rate α is small. Note that the performance of the stochastic algorithms does not depend on the estimation rate α .
- Hastily incorporating second-order information into gradient update methods may only bring small improvements in \mathcal{E} . When minimizing \mathcal{E} , all four algorithms are dominated by the polynomial factor in $1/\epsilon$. The second-order methods shown here only improve the logarithmic factors in $1/\epsilon$.
- Stochastic algorithms have the best generalization performance, despite being the worst optimization algorithms when viewed in the traditional sense. When minimizing \mathcal{E} , the batch algorithms have a dependence on α as well as the logarithmic factor in $1/\epsilon$.

For small-scale problems, \mathcal{E}_{opt} is negligible, so these conclusions do not necessarily apply.

One caveat is that stochastic algorithms are more vulnerable than batch algorithms to ill-conditioning. The results for both stochastic algorithms depend on the constant ν . The factor of κ in the results for GU becomes κ^2 in the results for SGU. As H_{ERM} becomes increasingly ill-conditioned, both κ and ν are liable to increase rapidly. One of the goals of decorrelating the inputs (Section 1.5) is to prevent H_{ERM} from having large eigenvalues [LeCun et al., 1998]. In light of the results from Table 3.2, the importance of decorrelating the inputs becomes even more apparent.

4 Optimization Algorithms for Neural Networks

4.1 Challenges in Training Neural Networks

Deep neural networks are used today to solve challenging problems in computer vision, natural language processing, and several other fields. Training these networks is prototypical of the large-scale problems discussed in the previous section. Our guidelines suggest that for such problems, optimization algorithms based on steepest descent should use mini-batches, and carefully evaluate the tradeoffs of incorporating second-order information. Almost all optimization algorithms used to train deep neural networks abide by these constraints.

First-order algorithms must effectively deal with the vanishing gradient problem and the issue of “pathological curvature” [Martens, 2010]. In Section 1.5, we mentioned that activation values in deeper layers of a network have a greater tendency to cluster around zero or very large values. Since the gradients in one layer depend multiplicatively on those in the next, the vanishing gradient problem usually affects the lowest levels most severely. We also saw that when features of the input are correlated, the weight updates

are no longer decoupled. Both of these phenomena suggest that *each* weight and bias in the network should be given its own learning rate [LeCun et al., 1998], something that is usually not done for traditional optimization problems.

The asymptotic rates for the stochastic algorithms in Table 3.2 suggest vulnerability to ill-conditioning of the Hessian, especially when second-order information is not used. The problem is in fact much more serious than this. Suppose that each weight and bias in the network is given its own learning rate. By viewing the full batch gradient update rule as a discrete time dynamical system, it can be shown that the optimal learning rate for the i th parameter scales the corresponding eigenvalue of the Hessian matrix to unity [LeCun et al., 1998]. In other words,

$$\eta_i^* := \frac{r}{\lambda_i},$$

where r is the size of the mini-batch. Intuitively, this means that we would like to amplify η_i along directions of low curvature, and decay η_i along directions of high curvature.

A small neural network has on the order of 10^6 parameters. This means that the gradient vector has 10^6 elements, while the Hessian matrix has 10^{12} elements. Even for a small network, computing the spectrum of the Hessian, let alone the Hessian itself, is completely intractible. LeCun et al. [1998] give several techniques for estimating the maximum eigenvalue of the Hessian matrix, but dividing all of the learning rates by this quantity would unnecessarily retard progress. This means that first-order algorithms must somehow deal with pathological curvature without access to second-order information.

Unsurprisingly, the most successful first-order optimization algorithms do this by making innovative use of first-order information. These techniques use a combination of the following three ingredients: momentum, a window of past updates, and the diagonal of the Hessian matrix (which can actually be computed efficiently). Despite the efficacy of these techniques, there is a small but definite gap in the performance between first-order algorithms and the two most successful second-order algorithms that have been adapted for use with neural networks: limited-memory BFGS (L-BFGS) and “Hessian free optimization” (HF) [Martens, 2010, Ngiam et al., 2011, Sutskever et al., 2013]. These second-order algorithms seem to make effective use of the curvature information of the Hessian matrix in order to explore regions of the parameter space that are inaccessible to first-order methods.

4.2 First-Order Optimization Algorithms

4.3 Second-Order Optimization Algorithms

References

Christopher M. Bishop. *Pattern Recognition and Machine Learning (Information Science and Statistics)*. Springer-Verlag New York, Inc., Secaucus, NJ, USA, 2006. ISBN 0387310738.

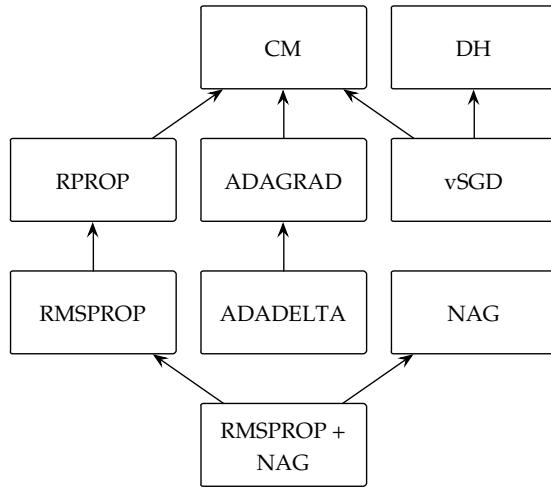


Figure 4.1: Hierarchy of the first-order optimization techniques and algorithms. An arrow from one block to another means that the algorithm represented by the first block uses or is based on the technique or algorithm represented by the second. Here, CM stands for “classical Nesterov momentum”, NAG for “Nesterov accelerated momentum”, and DH for “diagonal Hessian”.

Olivier Bousquet and Léon Bottou. The tradeoffs of large scale learning. In *Advances in neural information processing systems*, pages 161–168, 2008.

Xavier Glorot and Yoshua Bengio. Understanding the difficulty of training deep feedforward neural networks. In *International Conference on Artificial Intelligence and Statistics*, pages 249–256, 2010.

Alex Krizhevsky, Ilya Sutskever, and Geoffrey E Hinton. Imagenet classification with deep convolutional neural networks. In *Advances in neural information processing systems*, pages 1097–1105, 2012.

Y. LeCun, L. Bottou, G. Orr, and K. Muller. Efficient backprop. In G. Orr and Muller K., editors, *Neural Networks: Tricks of the trade*. Springer, 1998.

Y. LeCun, L. Bottou, Y. Bengio, and P. Haffner. Gradient-based learning applied to document recognition. In *Intelligent Signal Processing*, pages 306–351. IEEE Press, 2001.

James Martens. Deep learning via hessian-free optimization. In *Proceedings of the 27th International Conference on Machine Learning (ICML-10)*, pages 735–742, 2010.

Mehryar Mohri, Afshin Rostamizadeh, and Ameet Talwalkar. *Foundations of Machine Learning*. The MIT Press, 2012. ISBN 026201825X, 9780262018258.

Vinod Nair and Geoffrey E Hinton. Rectified linear units improve restricted boltzmann machines. In *Proceedings of the 27th International Conference on Machine Learning (ICML-10)*, pages 807–814, 2010.

Jiquan Ngiam, Adam Coates, Ahbik Lahiri, Bobby Prochnow, Quoc V Le, and Andrew Y Ng. On optimization methods for deep learning. In *Proceedings of the 28th International Conference on Machine Learning (ICML-11)*, pages 265–272, 2011.

Ilya Sutskever, James Martens, George Dahl, and Geoffrey Hinton. On the importance of initialization and momentum in deep learning. In *Proceedings of the 30th International Conference on Machine Learning (ICML-13)*, pages 1139–1147, 2013.

Matthew D Zeiler and Rob Fergus. Visualizing and understanding convolutional networks. In *Computer Vision–ECCV 2014*, pages 818–833. Springer, 2014.